

Jiayi Tian

+1 (805) 245 0298 | jiayi_tian@ucsb.edu | linkedin.com/in/jiayi-tian-32b9652a5/

Focus on Efficient LLM Training & Inference, Efficient CoT Reasoning via Algorithm-System Co-Design.

EDUCATION

University of California, Santa Barbara, <i>Ph.D. in Computer Engineering</i> CA, USA 3.93/4.0	Fall 2023 - ongoing
University of California, Santa Barbara, <i>M.S. in Computer Engineering</i> CA, USA 3.93/4.0	Fall 2023 - Fall 2025
Nanjing University, <i>B.Eng. in Electrical Engineering</i> China 4.50/5.0	Fall 2019 - Fall 2023

INDUSTRIAL EXPERIENCE

Intel Corporation, *Research Intern 2025* | Hillsboro, OR | **Mentor:** Souvik, Kundu Jun 2025 – Sep 2025

- Proposed SkipKV, a training-free KV-cache compression framework featuring sentence-level selective eviction and dynamic thoughts-type generation control for efficient chain-of-thought reasoning in multi-batch serving scenarios.
- Designed a semantic similarity metric to evict redundant sentences from the cache, improving reasoning coherence.
- Introduced a dynamic latent thoughts-type steering mechanism to enable concise and stable inference.
- Demonstrated strong results on long-reasoning tasks with LLMs: up to 26.7% higher accuracy, 1.6× shorter generation, and 1.7× higher throughput vs. SOTA under equal compression.
- Resulting paper accepted to MLSys, 2026.

Intel Corporation, *Research Intern 2024* | Hillsboro, OR | **Mentor:** Hai Li Jun 2024 - Sep 2024

- Proposed a tensor-compressed Transformer training accelerator on FPGA, optimizing compute ordering, dataflow, and memory allocation for LLMs.
- Designed a bidirectional tensor contraction scheme enabling substantial reduction in intermediate memory and compute cost during long-sequence training and inference.
- Built an HLS-based training engine achieving up to 51× memory efficiency and 4× energy efficiency compared with an Nvidia RTX 3090 GPU.
- Resulting paper accepted to IEEE TCAD, 2025.

AMD-Xilinx Technology, *Co-Op/Intern* | Beijing, China Jun 2023 - Sep 2023

- Developed a C++/HLS Transformer training framework with custom tensorized linear layers and nonlinear operations for LLM acceleration, achieved 30× ~ 52× saving in model size for end-to-end Transformer training.

SKILLS & RESEARCH INTERESTS

Languages & Tools Python, PyTorch, Huggingface, vLLM, C/C++

ML & NLP	Efficient Large Language Models (LLMs) Training/Inference, Efficient Large Reasoning Models (LRMs) (Model Compression, KV Cache Compression, Pruning, Low-rank decomposition, Early Exit, Knowledge Distillation, Quantization, Speculative Decoding/Reasoning), Efficient Agentic System
----------	---

SELECTED PUBLICATIONS & PREPRINTS

- RankGuide: Tensor-Rank-Guided Routing and Steering for Efficient Reasoning** [paper]
Jiayi Tian, Yupeng Su, Ryan Solgi, Souvik Kundu, Zheng Zhang, accepted to Third Conference on Language Modeling (COLM), 2026.
- SkipKV: Selective Skipping of KV Generation and Storage for Efficient Inference with Large Reasoning Models** [paper]
[code]
Jiayi Tian, Seyedarmin Azizi, Yequan Zhao, Erfan Baghaei Potraghloo, Sean McPherson, Sharath Nittur Sridhar, Zhengyang Wang, Zheng Zhang, Massoud Pedram, Souvik Kundu, accepted to MLSys, 2026.
- FLAT-LLM: Fine-grained Low-rank Activation Space Transformation for Large Language Model Compression** [paper]
[code]
Jiayi Tian, Ryan Solgi, Jinming Lu, Yifan Yang, Hai Li, Zheng Zhang, Findings of EACL, 2026.
- Ultra Memory-Efficient On-FPGA Training of Transformers via Tensor-Compressed Optimization** [paper]
Jiayi Tian, Jinming Lu, Hai Li, Xiangwei Wang, Cong (Callie) Hao, Ian Young, Zheng Zhang, IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems (TCAD), 2025.
- BEBERT: Efficient and robust binary ensemble BERT** [paper] [code]
Jiayi, Tian, Chao Fang, Haonan Wang, and Zhongfeng Wang, IEEE International Conference on Acoustics, Speech and Signal Processing (ICASSP), 2023.

Agentic System Acceleration via Speculation

Apr 2026 – ongoing

- Developed speculative tool-calling baselines for human-agent or multi-agent interaction benchmarks.
- Designing an agentic system with speculation, model routing, and meta-tool use to improve inference efficiency.

Tensor-Rank-Guided Steering and Routing for High-performance SRM Reasoning

Jan 2026. – Mar 2026.

- Proposed RankGuide, a framework that leverages tensor-rank signals from hidden states to accelerate LRMs reasoning, which can yield up to $1.75\times$ and $1.36\times$ latency benefit compared to LRM and SoTA collaborative inference framework, while maintaining or improving the accuracy.
- Designed a tensor-rank scoring metric on step-level hidden states to detect low-quality reasoning steps and selectively route them to larger models, improving the accuracy-latency trade-off.
- Developed a rank-based calibration pipeline that filters low-rank samples to construct high-quality steering vectors for inference-time hidden-state intervention, encouraging concise and stable reasoning.

Structural Pruning for Efficient LLM Inference via Low-rank Decomposition

Aug 2024 - May 2025

- Developed FLAT-LLM, a training-free, fine-grained compression method that leverages the low-rank structure of the activation space to transform and compress the model weights.
- Introduced a novel training-free rank selection algorithm that allocates ranks using a greedy redistribution strategy and can be integrated with existing low-rank LLM compression pipelines.
- Achieved strong performance on LLaMA-2, 3 and Mistral models with minimal calibration overhead (within minutes), validated across language modeling and downstream tasks.

Binary-Quantized Ensemble LLM for Fast and Robust Language Model Inference

Apr 2021 - Jun 2023

- Developed BEBERT, a novel quantization-ensemble strategy enabling efficient and accurate 1-bit BERT inference.
- Leveraged efficient knowledge distillation strategy for high training efficiency.
- Achieved $13\times$ model size reduction and $15\times$ compute savings over standard BERT with minimal accuracy loss.
- Proposed early-exit inference variant, further cutting compute by $20\% \sim 40\%$ on GLUE benchmark.